

# SpencerHann\_finalPresentation

March 19, 2020

## 1 Convolutional Neural Network

Spencer Hann EE 584 - Final Project

```
[1]: import numpy as np

from matplotlib import pyplot as plt
from cnn.data import preprocess_data

[2]: training_examples, training_targets = \
    preprocess_data("data/mnist_train.csv", max_rows=4000)
testing_examples, testing_targets = \
    preprocess_data("data/mnist_test.csv", max_rows=1000)

data = (training_examples, training_targets, testing_examples, testing_targets)
train_set = (training_examples, training_targets)
test_set = (testing_examples, testing_targets)
```

loading from file... done.

loading from file... done.

### 1.1 Proof of Concept: Simple Neural Net Layer

```
[3]: from cnn.cnn_basic import CNN, Layer, DenseSoftmaxLayer

[4]: class DenseLayer:
    def __init__(self, insize, outsize=10):
        self.insize = insize
        self.outsize = outsize

        self.w = np.random.randn(insize, outsize) / insize
        self.b = np.random.randn(outsize) / outsize

    def forward(self, image):
        image = image.flatten()
```

```

    # fully-connected/matmul phase
    result = np.dot(image, self.w) + self.b
    return result

```

```

[5]: layer1 = DenseLayer(28*28, 64)
     layer2 = DenseLayer(64, 10)

```

```

[6]: def forward(image, label):
     middle = layer1.forward(image)
     out = layer2.forward(middle)

     is_correct = np.argmax(out) == label
     return None, is_correct

```

```

[7]: def test(images, labels):
     n_correct = 0.0
     for image, label in zip(images, labels):
         _, c = forward(image, label)
         n_correct += c
     return n_correct / len(images)

```

```

[8]: accuracy = 100 * test(*test_set) # should be about 1 / n_classes
     print(f"Accuracy: {round(accuracy)}%")

```

Accuracy: 11.0%

## 1.2 Adding Back Propagation

```

[9]: class DenseLayer:
     def __init__(self, insize, outsize=10):
         self.insize = insize
         self.outsize = outsize

         self.w = np.random.randn(insize, outsize) / insize
         self.b = np.random.randn(outsize) / outsize

     def forward(self, image):
         self.last_image = image          # <<----
         image = image.flatten()

         # fully-connected/matmul phase
         fc = np.dot(image, self.w) + self.b
         self.last_fc = fc                 # <<----

         return fc

```

```

def backprop(self, loss_grad, lr=0.002):
    # output gradients wrt input, biases, weights
    ograd_input = self.w
    ograd_biases = 1
    ograd_weights = self.last_image.flatten()

    # loss gradients wrt input, biases, weights
    lgrad_input = ograd_input @ loss_grad
    lgrad_biases = ograd_biases * loss_grad
    lgrad_weights = ograd_weights[:, np.newaxis] @ loss_grad[np.newaxis]

    # update layer
    self.w += lr * lgrad_weights
    self.b += lr * lgrad_biases
    return lgrad_input.reshape(self.last_image.shape)

```

```

[10]: layer1 = DenseLayer(28*28, 64)
      layer2 = DenseLayer(64, 10)
      layers = (layer1, layer2,)

```

```

[11]: def forward(image, label):
      out = image
      for layer in layers:
          out = layer.forward(out)

      is_correct = np.argmax(out) == label
      loss = -np.log(out[label])          # <<-----

      return out, loss, is_correct

```

```

[12]: def learn(image, label):
      out, loss, correct = forward(image, label)

      if correct:
          return loss, correct

      # cross entropy gradient
      grad = np.zeros(10)
      grad[label] = - 1 / out[label]

      for layer in layers[::-1]:
          grad = layer.backprop(grad, lr=0.02)

      return loss, correct

```

```

[13]: def train(n_epochs, images, labels):
      n = len(images)

```

```

for epoch in range(n_epochs):
    ncorrect = 0
    for image, label in zip(images, labels):
        _, c = learn(image, label)
        ncorrect += c

    accuracy = round(100 * ncorrect / n)
    print(f"Epoch:{epoch}, Accuracy: {100 * ncorrect / n}")

```

### 1.3 Demonstration of Forward propogation

This is a randomly initialized network that we will test the feed forward functionality on. With 10 evenly represented output classes in our testing data, we expect to see roughly 10% accuracy.

```

[14]: from cnn.cnn_basic import (
        CNN,
        ConvolutionalLayer,
        MaxPoolingLayer,
        DenseSoftmaxLayer,
    )

```

```

[15]: cnn = CNN((
        ConvolutionalLayer(1,3,3, first_layer=True),
        MaxPoolingLayer(2),
        ConvolutionalLayer(3,1,3,),
        DenseSoftmaxLayer(14*14, 10),
    ))

```

```

[16]: # cnn.test(*test_set);

```

Approximately random performance, this is to be expected. It shows that feed forward is working properly.

### 1.4 Demonstration of Back Propagation

```

[17]: cnn = CNN(
        (
            ConvolutionalLayer(1,3,3, first_layer=True),
            DenseSoftmaxLayer(28*28*3, 10),
        ),
        lr = 0.01
    )

```

```

[18]: %time cnn.train_epochs(3, *train_set);

```

```
100%|      | 4000/4000 [02:46<00:00, 23.96it/s]
  0%|      | 2/4000 [00:00<03:21, 19.87it/s]

Epoch 0/3: 1.91 loss, 39.32% accurate

100%|      | 4000/4000 [02:46<00:00, 23.96it/s]
  0%|      | 2/4000 [00:00<03:38, 18.26it/s]

Epoch 1/3: 2.01 loss, 47.85% accurate

100%|      | 4000/4000 [03:20<00:00, 19.92it/s]

Epoch 2/3: 2.21 loss, 49.30% accurate
CPU times: user 18min 55s, sys: 12min 19s, total: 31min 15s
Wall time: 8min 54s
```

```
[19]: cnn.test(*test_set);
```

```
100%|      | 1000/1000 [00:14<00:00, 71.36it/s]

Test: 2.74 loss, 41.40% accurate
```

Though this is significantly better than random, it is still not as high as I'd like it to be. I believe the reason for this is that the network in its current state is relatively low-capacity, and does not contain any non-linearities.